

Authors' Instructions: Preparation of Camera-Ready Contributions to SCITEPRESS Proceedings

First Author Name¹^a, Second Author Name¹^b and Third Author Name²^c

¹*Institute of Problem Solving, XYZ University, My Street, MyTown, MyCountry*

²*Department of Computing, Main University, MySecondTown, MyCountry*

{first_author; second_author}@ips.xyz.edu, third_author@dc.mu.edu

Keywords: The paper must have at least one keyword. The text must be set to 9-point font size and without the use of bold or italic font style. For more than one keyword, please use a comma as a separator. Keywords must be titlecased.

Abstract: The abstract should summarize the contents of the paper and should contain at least 70 and at most 200 words. The text must be set to 9-point font size.

1 Introduction

>>> To be completed. . .

- We present “search coaching” as a way to instill knowledge faithfully to a learner;
- We compare this to our own and others’ previous works;
- We maybe discuss a bit about LLMs. (?)

2 Related literature

>>> To be completed. . .

- Related works as in the proxy coaching paper;
- A brief update on what is the current state of the art;
- Maybe some works that draw links to the planning aspect of this paper.

3 Experimental setup

In this Section we empirically explore search coaching, by having a machine learner iteratively seek advice by a coach, as in (Markos et al., 2022), with the aim of iteratively improving its search capacity, as per

the coach’s explicit search policy. Given a fixed goal state, g , the learner applies its search policy, P_{learner} , to identify a path, p , starting from an arbitrary start state s and ending at g . Upon lack of knowledge, the learner seeks advice from its coach, which responds according to its own policy, P_{coach} . A *coaching step* is a single learner request for advice and the corresponding coach response, while a *coaching session* comprises of the full learning cycle, where the learner eventually crafts one such path or the process is aborted.

3.1 Sorting coaching

In this series of coaching sessions we have implemented two variants of a machine coach, aiming to explain to an initial ignorant agent about how to sort a list of numbers, as a rough equivalent to linear preference elicitation. The first coach, C_{bubble} , uses bubble sort as its underlying policy while the second one, C_{quick} , uses quicksort to provide advice from. We explored different state sizes, n , ranging from 1 to 20, running $m = 100$ iterations for each value of n . We have also explored two different types of advice for each coach: (A1) *full advice*, where each piece of advice depends on the entire state, and; (A2) *partial advice*, where each piece of advice depends on the specific part of the state to be improved. Moreover, for each advice type we have considered three learner configurations regarding advice memory: (M1) *no memory* across different iterations for the same state size, n , i.e., coaching sessions are pairwise indepen-

Add a few words about how the learner improves its search capacity in the context of searching for a past sub-optimal state the piece of advice according to its own policy that improves a(n implicit / explicit) goal.

^a <https://orcid.org/0000-0000-0000-0000>

^b <https://orcid.org/0000-0000-0000-0000>

^c <https://orcid.org/0000-0000-0000-0000>

dent; (M2) *short memory* across different iterations for the same state size, n , i.e., coaching sessions across different values of n are independent but not within the same value for n , and; (M3) *long memory* across all values of state size, n , so all coaching sessions are correlated, building on knowledge from previous test cases. In all cases the learner seeks advice upon lack of knowledge for further searching, returning a partial path p starting from s , while the coach adopts one of the following *coaching styles*: (C1) in *reactive* coaching, denoted by a superscript a , the coach provides advice on learner’s last wrong decision in p ; (C2) in *proactive* coaching, denoted by a superscript p , the coach provides advice on learner’s first wrong decision in p ; (C3) in *reflexive* coaching, denoted by a superscript r , the coach provides advice at a random wrong decision in p .

In Figures 1 we present the number of coaching interactions required to reach the goal state in each case as a function of the state size, n . Regarding the full advice condition (Fig. 1a), the utilization of either short or long term memory does not appear to have any particular effect on the number of coaching interactions, which was expected given the overspecific nature of provided advice, being fully tailored to the state it triggered the learner to seek for that piece of advice. On the contrary, partial advice (Fig. 1b) offer a significant advantage to the learner, in both short and long memory conditions. As for the short memory case,

Data: s, g , learner, coach

Result: learner with updated search policy.

$advice \leftarrow \text{null}$;

do

$learner.update_hypothesis(advice)$;
 $path \leftarrow learner.search_path(s, g)$;
 $advice \leftarrow coach.evaluate(path, s, g)$;

while $advice$ is not null;

Algorithm 1: Coaching process. s, g denote the start and goal states respectively.

3.2 Clustering Coaching

- We explore clustering in various settings, using an inertia-based algorithm.
- The setting is **similar** in terms of the coaching process, but there do exist **domain differences**:
 - Algorithms rely mostly on full– and not partial–state advice;
 - We are exploring how different levels of “generalization freedom” affect learning (with memory).

4 Results

>>> To be completed...


- Find a metric to quantify learner’s fidelity (brainstorming follows):
 - Maybe the number of advice itself is an indicator, at least in terms of efficacy – it still might be the case that the learner achieves the same result through another path;
 - Different coaching styles (reactive, reflexive, proactive) should affect faithfulness / fidelity, so this metric should capture it.
 - The absolute (or relative?) frequency of differences between the learner’s and the coach’s paths on a given benchmark set could be the most fine grained metric. So, for instance, if the learner goes $s \rightarrow a \rightarrow b \rightarrow c \rightarrow g$ while the coach would go $s \rightarrow a \rightarrow b \rightarrow d \rightarrow g$, we have an absolute (relative) difference of 1 (0.33, counting only intermediate nodes).
 - In the above, we might also count the goal node in the relative frequency case. Also, we should somehow measure distances among paths of different length.

5 Conclusions and future work


>>> To be completed...

REFERENCES


- Markos, V. T., Thoma, M., and Michael, L. (2022). Machine coaching with proxy coaches. In Kuhlmann, I., Mumford, J., and Sarkadi, S., editors, *Proceedings of the 1st Workshop on Argumentation & Machine Learning co-located with 9th International Conference on Computational Models of Argument (COMMA 2022)*, Cardiff, Wales, September 13th, 2022, volume 3208 of *CEUR Workshop Proceedings*, pages 45–64. CEUR-WS.org.




(a) Full condition advice. reactive coaching




(b) Partial condition advice. reactive coaching




(c) Full condition advice. proactive coaching



(d) Partial condition advice. proactive coaching



(e) Full condition advice, reflexive coaching



(f) Partial condition advice, reflexive coaching

Figure 1: >>> To be completed...